

# Logistic Regression Churn Prediction

June 8, 2022

## 1 Logistic Regression Churn Prediction Model

### 1.1 Importing Libraries

```
[1]: from sklearn import linear_model
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
import scipy
from scipy import stats
import matplotlib.pyplot as plt
import plotly.express as px
```

### 1.2 Data Inspection & Cleaning

```
[2]: #Importing the dataset and looking at the attributes of the dataset
```

```
df = pd.read_csv('data_regression.csv')
df.columns
```

```
[2]: Index(['year', 'customer_id', 'phone_no', 'gender', 'age',
          'no_of_days_subscribed', 'multi_screen', 'mail_subscribed',
          'weekly_mins_watched', 'minimum_daily_mins', 'maximum_daily_mins',
          'weekly_max_night_mins', 'videos_watched', 'maximum_days_inactive',
          'customer_support_calls', 'churn'],
          dtype='object')
```

```
[3]: df.head(10)
```

```
[3]:   year  customer_id  phone_no  gender  age  no_of_days_subscribed  \
0  2015         100198  409-8743  Female   36                62
1  2015         100643  340-5930  Female   39                149
2  2015         100756  372-3750  Female   65                126
3  2015         101595  331-4902  Female   24                131
4  2015         101653  351-8398  Female   40                191
```

5	2015	101953	329-6603	NaN	31	65
6	2015	103051	416-1845	NaN	54	59
7	2015	103225	348-7193	Female	40	50
8	2015	103408	413-4039	Male	61	205
9	2015	103676	338-5207	Male	31	63

	multi_screen	mail_subscribed	weekly_mins_watched	minimum_daily_mins	\
0	no	no	148.35	12.2	
1	no	no	294.45	7.7	
2	no	no	87.30	11.9	
3	no	yes	321.30	9.5	
4	no	no	243.00	10.9	
5	no	no	193.65	12.7	
6	no	no	239.25	10.2	
7	no	no	196.65	5.6	
8	no	yes	263.70	7.8	
9	no	no	316.80	12.3	

	maximum_daily_mins	weekly_max_night_mins	videos_watched	\
0	16.81	82	1	
1	33.37	87	3	
2	9.89	91	1	
3	36.41	102	4	
4	27.54	83	7	
5	21.95	111	6	
6	27.12	106	4	
7	22.29	88	9	
8	29.89	64	5	
9	35.90	58	2	

	maximum_days_inactive	customer_support_calls	churn
0	4.0	1	0.0
1	3.0	2	0.0
2	4.0	5	1.0
3	3.0	3	0.0
4	3.0	1	0.0
5	4.0	4	1.0
6	NaN	0	0.0
7	NaN	5	1.0
8	3.0	2	0.0
9	4.0	0	0.0

We can see that the gender and maximum\_days\_inactive is having the NAN values. Gender, screen, and mail are categorical variables. There are also other behavioural variables like weekly\_mins\_watched, minimum\_daily\_mins, maximum\_daily\_mins, weekly\_max\_night\_mins, videos\_watched, maximum\_days\_inactive, customer\_support\_calls, and churn.

```
[4]: #Defining Inspection and Cleaning function
```

```
def inspection(df):

    import pandas as pd
    import seaborn as sns

    print('Types of variables we are working with: ')
    print(df.dtypes, "\n")

    print('Total number of samples with missing values')
    print(df.isnull().any(axis=1).sum(), "\n")

    print('Total missing variables per attribute/column')
    print(df.isnull().sum(), "\n")

    print('Map of missing values')
    sns.heatmap(df.isnull())
```

```
[5]: inspection(df)
```

Types of variables we are working with:

year	int64
customer_id	int64
phone_no	object
gender	object
age	int64
no_of_days_subscribed	int64
multi_screen	object
mail_subscribed	object
weekly_mins_watched	float64
minimum_daily_mins	float64
maximum_daily_mins	float64
weekly_max_night_mins	int64
videos_watched	int64
maximum_days_inactive	float64
customer_support_calls	int64
churn	float64
dtype:	object

Total number of samples with missing values  
82

Total missing variables per attribute/column


year	0
customer_id	0
phone_no	0

```

gender                24
age                   0
no_of_days_subscribed 0
multi_screen          0
mail_subscribed       0
weekly_mins_watched   0
minimum_daily_mins    0
maximum_daily_mins    0
weekly_max_night_mins 0
videos_watched        0
maximum_days_inactive 28
customer_support_calls 0
churn                 35
dtype: int64

```

Map of missing values



From the heatmap, it doesn't seem like there is a lot of missing data. There are some values

for Churn coloumn which are missing. Mostly, those are in the first 1000 samples. Few values of gneder and maximum\_days\_inactive are also missing. These can be seen in first half of the data. There is not much missing data, but again we need to get rid of the data which is holding null values. So, we will remove all the rows with missing values from the data frame.

[6]: *#Using dropna to remove missing values*

```
df = df.dropna()
df
```

```
[6]:      year  customer_id  phone_no  gender  age  no_of_days_subscribed  \
0    2015      100198  409-8743  Female   36                62
1    2015      100643  340-5930  Female   39                149
2    2015      100756  372-3750  Female   65                126
3    2015      101595  331-4902  Female   24                131
4    2015      101653  351-8398  Female   40                191
...    ...    ...    ...    ...    ...    ...
1990  2015      993714  364-1969   Male   32                61
1991  2015      993815  387-5891   Male   49                50
1992  2015      994954  329-3222  Female   42                119
1996  2015      998086  383-9255   Male   45                127
1999  2015      999961  414-1496   Male   37                73

      multi_screen  mail_subscribed  weekly_mins_watched  minimum_daily_mins  \
0              no              no              148.35              12.2
1              no              no              294.45              7.7
2              no              no              87.30              11.9
3              no              yes              321.30              9.5
4              no              no              243.00              10.9
...    ...    ...    ...    ...    ...
1990              no              no              67.50              9.8
1991              yes              yes              460.65              8.0
1992              no              yes              176.70              7.6
1996              no              no              273.45              9.3
1999              no              no              326.70              10.3

      maximum_daily_mins  weekly_max_night_mins  videos_watched  \
0              16.81              82              1
1              33.37              87              3
2              9.89              91              1
3              36.41              102             4
4              27.54              83              7
...    ...    ...    ...
1990              7.65              94              6
1991              52.21             109              3
1992              20.03              76              3
1996              30.99             116              3
```

1999	37.03	89	6
------	-------	----	---

	maximum_days_inactive	customer_support_calls	churn
0	4.0	1	0.0
1	3.0	2	0.0
2	4.0	5	1.0
3	3.0	3	0.0
4	3.0	1	0.0
...	...	...	...
1990	3.0	2	0.0
1991	3.0	0	0.0
1992	3.0	3	0.0
1996	3.0	1	0.0
1999	3.0	1	1.0

[1918 rows x 16 columns]

```
[7]: df.describe()
```

```
[7]:
```

	year	customer_id	age	no_of_days_subscribed	\
count	1918.0	1918.000000	1918.000000	1918.000000	
mean	2015.0	562267.639208	38.659020	100.065693	
std	0.0	257172.549309	10.150896	39.923242	
min	2015.0	100198.000000	18.000000	1.000000	
25%	2015.0	341491.750000	32.000000	73.000000	
50%	2015.0	579594.500000	37.000000	99.000000	
75%	2015.0	778377.250000	43.000000	127.000000	
max	2015.0	999961.000000	82.000000	243.000000	

	weekly_mins_watched	minimum_daily_mins	maximum_daily_mins	\
count	1918.000000	1918.000000	1918.000000	
mean	269.919291	10.180553	30.591413	
std	80.541705	2.771199	9.128036	
min	0.000000	0.000000	0.000000	
25%	218.587500	8.400000	24.775000	
50%	269.550000	10.200000	30.550000	
75%	324.000000	12.000000	36.720000	
max	526.200000	20.000000	59.640000	

	weekly_max_night_mins	videos_watched	maximum_days_inactive	\
count	1918.000000	1918.000000	1918.000000	
mean	100.400938	4.484880	3.247132	
std	19.569822	2.477764	0.805840	
min	42.000000	0.000000	0.000000	
25%	87.000000	3.000000	3.000000	
50%	101.000000	4.000000	3.000000	
75%	114.000000	6.000000	4.000000	

max	175.000000	19.000000	6.000000
-----	------------	-----------	----------

	customer_support_calls	churn
count	1918.000000	1918.000000
mean	1.539625	0.131908
std	1.300553	0.338479
min	0.000000	0.000000
25%	1.000000	0.000000
50%	1.000000	0.000000
75%	2.000000	0.000000
max	9.000000	1.000000

Attribute year is having a zero standard deviation because all the elements of attribute year are same, i.e. the year 2015. The dataset contains only the data about the customers who were there on the platform during the year 2015. There are 1918 rows, i.e. 1918 customers in total. This is after we remove the null values from the dataset. Age of the customer varies from 18 to 82. The main attribute used to describe if the customer churned or not is based on maximum days the customer was inactive. The minimum value for this attribute is 0, which is reasonable as some of the customers may not be inactive on the platform, but the maximum value is 6, i.e. some of the customers were inactive for longer time and result into getting churned. This description says a lot about the customers, i.e. the average customers were of age 38-39 and on an average 0.13% of customers are getting churned.

```
[8]: #Printing the correlation matrix
corr = df.corr()
corr
```

```
[8]:
```

	year	customer_id	age	no_of_days_subscribed	\
year	NaN	NaN	NaN	NaN	
customer_id	NaN	1.000000	0.021881	-0.019180	
age	NaN	0.021881	1.000000	0.035426	
no_of_days_subscribed	NaN	-0.019180	0.035426	1.000000	
weekly_mins_watched	NaN	-0.010410	0.019586	-0.002089	
minimum_daily_mins	NaN	0.040254	-0.008557	0.015247	
maximum_daily_mins	NaN	-0.010415	0.019598	-0.002095	
weekly_max_night_mins	NaN	0.000648	0.015150	0.001290	
videos_watched	NaN	0.061229	-0.003876	0.012856	
maximum_days_inactive	NaN	0.047551	0.001507	0.017720	
customer_support_calls	NaN	-0.034940	-0.002848	0.011272	
churn	NaN	-0.054260	0.015982	0.009627	

	weekly_mins_watched	minimum_daily_mins	\
year	NaN	NaN	
customer_id	-0.010410	0.040254	
age	0.019586	-0.008557	
no_of_days_subscribed	-0.002089	0.015247	
weekly_mins_watched	1.000000	-0.015180	

minimum_daily_mins	-0.015180	1.000000
maximum_daily_mins	1.000000	-0.015178
weekly_max_night_mins	0.039856	0.011446
videos_watched	0.027869	0.046493
maximum_days_inactive	-0.012410	0.931296
customer_support_calls	-0.031239	-0.003817
churn	0.165991	0.072308

	maximum_daily_mins	weekly_max_night_mins	\
year	NaN	NaN	
customer_id	-0.010415	0.000648	
age	0.019598	0.015150	
no_of_days_subscribed	-0.002095	0.001290	
weekly_mins_watched	1.000000	0.039856	
minimum_daily_mins	-0.015178	0.011446	
maximum_daily_mins	1.000000	0.039858	
weekly_max_night_mins	0.039858	1.000000	
videos_watched	0.027870	-0.003355	
maximum_days_inactive	-0.012410	0.032647	
customer_support_calls	-0.031250	-0.013568	
churn	0.165989	0.006029	

	videos_watched	maximum_days_inactive	\
year	NaN	NaN	
customer_id	0.061229	0.047551	
age	-0.003876	0.001507	
no_of_days_subscribed	0.012856	0.017720	
weekly_mins_watched	0.027869	-0.012410	
minimum_daily_mins	0.046493	0.931296	
maximum_daily_mins	0.027870	-0.012410	
weekly_max_night_mins	-0.003355	0.032647	
videos_watched	1.000000	0.026169	
maximum_days_inactive	0.026169	1.000000	
customer_support_calls	-0.009850	-0.006362	
churn	-0.018457	0.046809	

	customer_support_calls	churn
year	NaN	NaN
customer_id	-0.034940	-0.054260
age	-0.002848	0.015982
no_of_days_subscribed	0.011272	0.009627
weekly_mins_watched	-0.031239	0.165991
minimum_daily_mins	-0.003817	0.072308
maximum_daily_mins	-0.031250	0.165989
weekly_max_night_mins	-0.013568	0.006029
videos_watched	-0.009850	-0.018457
maximum_days_inactive	-0.006362	0.046809



customer_support_calls	1.000000	0.212678
churn	0.212678	1.000000

### 1.3 Encoding Categorical Variables

```
[9]: df.head(5)
```

```
[9]:   year  customer_id  phone_no  gender  age  no_of_days_subscribed  \
0  2015      100198  409-8743  Female   36                62
1  2015      100643  340-5930  Female   39                149
2  2015      100756  372-3750  Female   65                126
3  2015      101595  331-4902  Female   24                131
4  2015      101653  351-8398  Female   40                191

   multi_screen  mail_subscribed  weekly_mins_watched  minimum_daily_mins  \
0           no                no                148.35                12.2
1           no                no                294.45                7.7
2           no                no                87.30                11.9
3           no                yes                321.30                9.5
4           no                no                243.00                10.9

   maximum_daily_mins  weekly_max_night_mins  videos_watched  \
0                16.81                    82                1
1                33.37                    87                3
2                 9.89                    91                1
3                36.41                   102                4
4                27.54                    83                7

   maximum_days_inactive  customer_support_calls  churn
0                 4.0                1      0.0
1                 3.0                2      0.0
2                 4.0                5      1.0
3                 3.0                3      0.0
4                 3.0                1      0.0
```

We can see our dataset has 3 categorical variables, i.e. gender, multi\_screen, and mail\_subscribed. We can look at the unique values of each attribute as follows

```
[10]: df.gender.unique()
```

```
[10]: array(['Female', 'Male'], dtype=object)
```

```
[11]: df.multi_screen.unique()
```

```
[11]: array(['no', 'yes'], dtype=object)
```

```
[12]: df.mail_subscribed.unique()
```

```
[12]: array(['no', 'yes'], dtype=object)
```

```
[13]: #Function for encoding categorical variables using scikit learn function called as ordinal encoder
```

```
def enc_cat (df, variables):  
    from sklearn.preprocessing import OrdinalEncoder  
    ord_enc = OrdinalEncoder()  
  
    for v in variables:  
        name = v+'_code'  
        df[name] = ord_enc.fit_transform(df[[v]])  
        print('The encoded values for ' + v + ' are: ' )  
        print(df[name].unique())
```

```
[14]: enc_cat(df, ['gender', 'multi_screen', 'mail_subscribed'])
```

The encoded values for gender are:

```
[0. 1.]
```

The encoded values for multi\_screen are:

```
[0. 1.]
```

The encoded values for mail\_subscribed are:

```
[0. 1.]
```

```
[15]: df
```

```
[15]:
```

	year	customer_id	phone_no	gender	age	no_of_days_subscribed	\
0	2015	100198	409-8743	Female	36	62	
1	2015	100643	340-5930	Female	39	149	
2	2015	100756	372-3750	Female	65	126	
3	2015	101595	331-4902	Female	24	131	
4	2015	101653	351-8398	Female	40	191	
...	...	...	...	...	...	...	
1990	2015	993714	364-1969	Male	32	61	
1991	2015	993815	387-5891	Male	49	50	
1992	2015	994954	329-3222	Female	42	119	
1996	2015	998086	383-9255	Male	45	127	
1999	2015	999961	414-1496	Male	37	73	

	multi_screen	mail_subscribed	weekly_mins_watched	minimum_daily_mins	\
0	no	no	148.35	12.2	
1	no	no	294.45	7.7	
2	no	no	87.30	11.9	
3	no	yes	321.30	9.5	
4	no	no	243.00	10.9	
...	...	...	...	...	
1990	no	no	67.50	9.8	

1991	yes	yes	460.65	8.0
1992	no	yes	176.70	7.6
1996	no	no	273.45	9.3
1999	no	no	326.70	10.3

	maximum_daily_mins	weekly_max_night_mins	videos_watched	\
0	16.81	82	1	
1	33.37	87	3	
2	9.89	91	1	
3	36.41	102	4	
4	27.54	83	7	
...	...	...	...	
1990	7.65	94	6	
1991	52.21	109	3	
1992	20.03	76	3	
1996	30.99	116	3	
1999	37.03	89	6	

	maximum_days_inactive	customer_support_calls	churn	gender_code	\
0	4.0	1	0.0	0.0	
1	3.0	2	0.0	0.0	
2	4.0	5	1.0	0.0	
3	3.0	3	0.0	0.0	
4	3.0	1	0.0	0.0	
...	...	...	...	...	
1990	3.0	2	0.0	1.0	
1991	3.0	0	0.0	1.0	
1992	3.0	3	0.0	0.0	
1996	3.0	1	0.0	1.0	
1999	3.0	1	1.0	1.0	

	multi_screen_code	mail_subscribed_code
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	1.0
4	0.0	0.0
...	...	...
1990	0.0	0.0
1991	1.0	1.0
1992	0.0	1.0
1996	0.0	0.0
1999	0.0	0.0

[1918 rows x 19 columns]

We can see that three new variables are created at the end of the dataframe called as gender\_code,

multi\_screen\_code, and mail\_subscribed\_code. All of those have only two types of values, i.e. 0 & 1.

## 1.4 Exploratory Data Analysis

```
[16]: #Function to return scatterplots of all the variables in the dataset against
      ↪ the classification variable

def plot1(df, cols_to_exclude, class_col):

    import numpy as np
    import seaborn as sns
    import warnings

    warnings.filterwarnings('ignore')

    #cleaning up of non-numerical columns
    cols = df.select_dtypes(include=np.number).columns.tolist() #finding all
    ↪ the numerical cols from df
    X = df[cols] #creating df
    ↪ only with numerical cols
    X = X[X.columns.difference(cols_to_exclude)] #removing
    ↪ columns to exclude

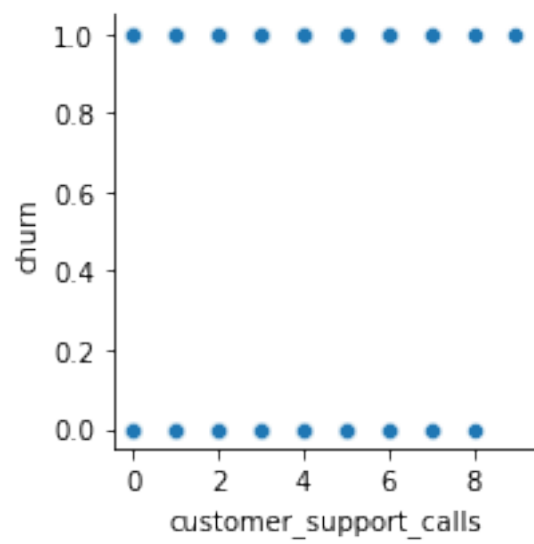

    #function dataframe.columns.difference give the complement of the values
    ↪ that we provide as argument.
    #here we are providing the cols to be excluded list as arg, so it will
    ↪ return all other cols other than those

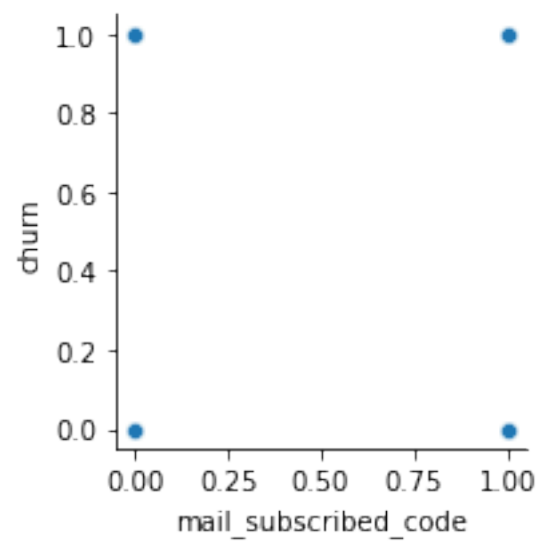

    for col in X.columns.difference([class_col]): #selecting all
    ↪ cols in list except the churn col
        g = sns.FacetGrid(df)

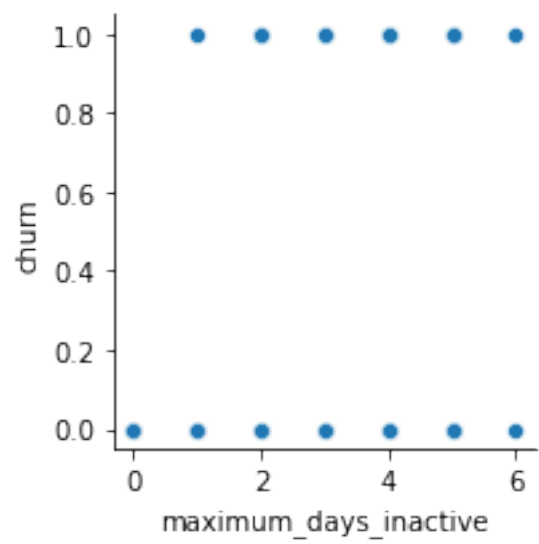

        #Facetgrid maps dataset onto multiple axes arrayed in a grid of rows or
    ↪ columns that correspond to levels of variables
        #in the dataset. Here we have only two variables compared at a time, so
    ↪ it is not using hue.

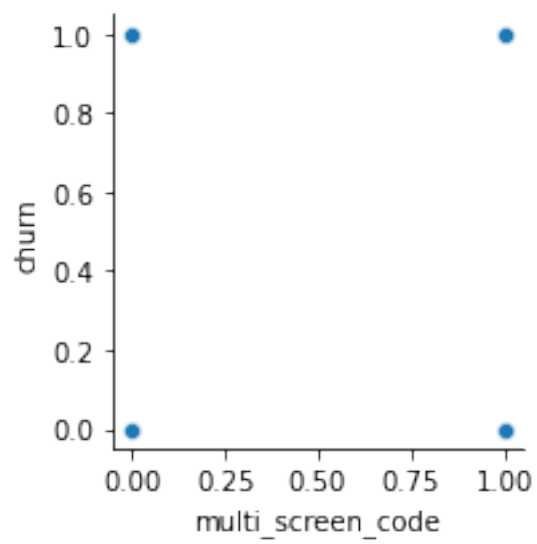

        #Applying a plot function to each facet's subset of the data
        g.map(sns.scatterplot, col, class_col)
```

```
[17]: plot1(df, ['customer_id', 'phone_no', 'year'], 'churn')
```

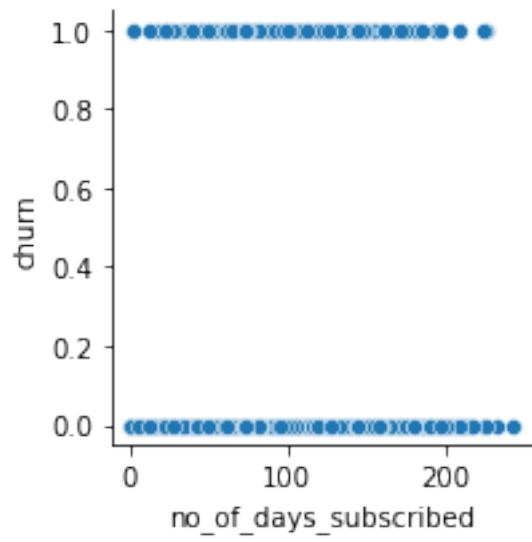



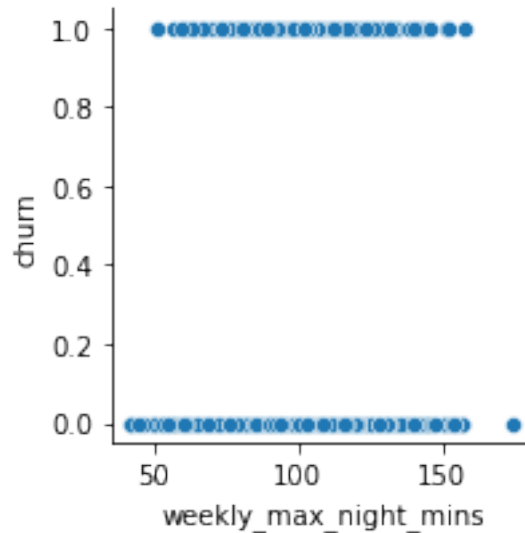












Unfortunately, it can be seen from the graph that there aren't a many obvious patterns when it comes to logistic regression. For example, if we see the above graph for churn vs weekly\_mins\_watched, we can see that people that churn and do not churn are all over the scale. Similarly, the graph for weekly\_max\_night\_mins, videos\_watched, no\_of\_days\_subscribed, minimum\_daily\_mins, etc have when compared with churn.

[18]: *#Function to return the pair-plots of classification variable versus all other variables*

```
def plot2(df, class_col, cols_to_exclude):
```

```

import numpy as np
import seaborn as sns

#cleaning up of non-numerical columns
cols = df.select_dtypes(include=np.number).columns.tolist() #finding all
↳ the numerical cols from df
X = df[cols] #creating df
↳ only with numerical cols
X = X[X.columns.difference(cols_to_exclude)] #removing
↳ columns to exclude

#function dataframe.columns.difference give the complement of the values
↳ that we provide as argument.
#here we are providing the cols to be excluded list as arg, so it will
↳ return all other cols other than those

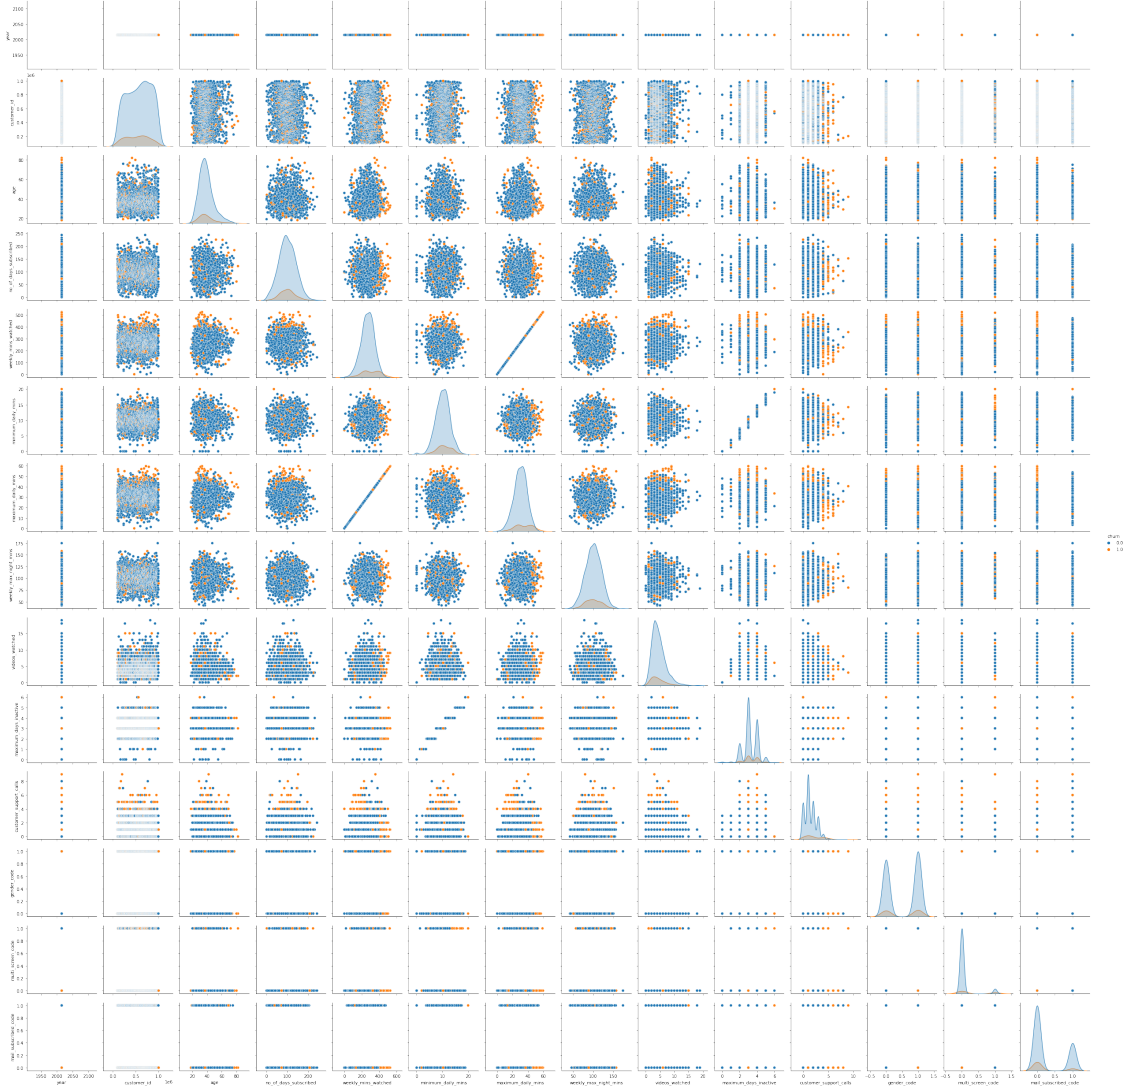
sns.pairplot(df, hue = class_col)

```

```

[19]: plot2(df, class_col='churn', cols_to_exclude=['customer_id', 'phone_no',
↳ 'year'])

```



This is the pair-plot of all the variables against each other. The diagonal graphs are all histograms. The one with orange dots are the customers who got churned and the blue ones were not churned.

For example, if we check the graph of `weeklymins_watched` versus `weeklymins_left`, we can see that more people got churned when `weeklymins_watched` went beyond a certain limit. This can be an explanation for the situation where people are coming on the platform to watch one particular series and then are getting churned.

Similarly, if we see the graph of `maximum_dailymins` and `weeklymins_watched` we can see that there is a linear relation which is growing increasingly. As the people are on the platform for few weeks their daily watch time is also increasing and most of them are getting churned when they are on the platform for long time or beyond some limit.

Also, when we see the relation between other variables and the encoded variables it will only show data points at values 0 and 1. This is because these variables were encoded with the values 0 and 1.

There is one more relation which is noticeable, i.e. as the minimum\_daily\_mins are increasing from 0 to 20 the maximum\_days\_inactive count is increasing from 0 to 6. Some people still stay on the platform and some of them get churned.

## 1.5 Running Logistic Regression

### 1.5.1 1. Running model using statsmodels.api

```
[22]: def logistic_regression(df, class_col, cols_to_exclude):

    #cleaning the dataframe for logistic regression with the columns which
    ↳we're not using, i.e. phone no, customer_id, & year
    cols = df.select_dtypes(include=np.number).columns.tolist()
    X = df[cols]
    X = X[X.columns.difference(cols_to_exclude)]
    X = X[X.columns.difference([class_col])]

    #We're running the logistic regression here using two ways:
    # 1. Using statsmodels.api library which bring the coeff with itself and is
    ↳easy to visualize
    # 2. Using sklearn for logistic regression

    # 1.
    import statsmodels.api as sm
    y = df[class_col]
    logit_1 = sm.Logit(y, X)
    result_1 = logit_1.fit()
    print(result_1.summary2())
```

```
[23]: logistic_regression(df, class_col = 'churn', cols_to_exclude=['customer_id',
    ↳'phone_no', 'year'])
```

Optimization terminated successfully.

Current function value: 0.336585

Iterations 7

Results: Logit

```
=====
Model:                Logit                Pseudo R-squared:    0.137
Dependent Variable:    churn                AIC:                1315.1404
Date:                 2022-06-08 18:53       BIC:                1381.8488
No. Observations:     1918                Log-Likelihood:     -645.57
Df Model:             11                  LL-Null:            -748.02
Df Residuals:         1906                LLR p-value:        7.1751e-38
Converged:            1.0000                Scale:              1.0000
No. Iterations:       7.0000

-----
                        Coef.  Std.Err.    z    P>|z|    [0.025    0.975]
-----
```

age	-0.0208	0.0068	-3.0739	0.0021	-0.0340	-0.0075
customer_support_calls	0.4246	0.0505	8.4030	0.0000	0.3256	0.5237
gender_code	-0.2144	0.1446	-1.4824	0.1382	-0.4979	0.0691
mail_subscribed_code	-0.7529	0.1798	-4.1873	0.0000	-1.1053	-0.4005
maximum_daily_mins	-3.7125	25.2522	-0.1470	0.8831	-53.2058	45.7809
maximum_days_inactive	-0.7870	0.2473	-3.1828	0.0015	-1.2716	-0.3024
minimum_daily_mins	0.2075	0.0727	2.8555	0.0043	0.0651	0.3499
multi_screen_code	1.9511	0.1831	10.6562	0.0000	1.5923	2.3100
no_of_days_subscribed	-0.0045	0.0018	-2.5572	0.0106	-0.0080	-0.0011
videos_watched	-0.0948	0.0317	-2.9954	0.0027	-0.1569	-0.0328
weekly_max_night_mins	-0.0169	0.0032	-5.3119	0.0000	-0.0231	-0.0107
weekly_mins_watched	0.4248	2.8619	0.1484	0.8820	-5.1844	6.0340

=====

We can see that the model iterated 7 times to reach the best optimization. Maximum number of iterations using Logit for statsmodels.api is 35, after this optimization fails. The value of log-likelihood is -645.57, i.e. we get our best logistic regression model at this value of LL. Also, the Psuedo R-squared value is very low, i.e. 13.7% which is not good and it says that our choice of model is not good.

According to the p-values, gender\_code, maximum\_daily\_mins, and weekly\_mins\_watched are not significant.

[24]: *#Interpreting coefficients of age*

```
import math
math.exp(-0.0208)
```

[24]: 0.9794148279480585

This means that an addition of year in customers age increases the odds of churning by 0.97. This is not a lot.

[25]: *#Interpreting coefficients of multi\_screen\_code*

```
import math
math.exp(1.9511)
```

[25]: 7.036423390843554

This shows that having a multi\_screen\_code changes the odd of churning by 7.03% and also the p-value of this attribute shows that it is very significant as it is 0.0000.

## 1.5.2 2. Running model using sklearn

[31]: `def prep_model(df, class_col, cols_to_exclude):`

```
    from sklearn.model_selection import train_test_split
```

```

import numpy as np

#cleaning the dataframe for logistic regression with the columns which
↪we're not using, i.e. phone no, customer_id, & year
cols = df.select_dtypes(include=np.number).columns.tolist()
X = df[cols]
X = X[X.columns.difference(cols_to_exclude)]
X = X[X.columns.difference([class_col])]

y = df[class_col]

#Declaring globally so that we can call this variables outside this
↪function also.
global X_train, X_test, y_train, y_test

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↪random_state = 0)

```

```

[32]: def running_model(X_train, X_test, y_train, y_test):

    from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import roc_auc_score, classification_report

    #Defining the logistic model globally to use it outside the function
    global logreg

    #Fitting the logistic regression model from sklearn
    logreg = LogisticRegression(random_state = 13)
    logreg.fit(X_train, y_train)

    #Predicting y values
    global y_pred #Defining globally to use outside the function

    y_pred = logreg.predict(X_test)

    logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))

    print(classification_report(y_test, y_pred))
    print('The area under the curve is: %0.2f'%logit_roc_auc)

```

## 1.6 Model Evaluation

```

[33]: prep_model(df, class_col='churn', cols_to_exclude=['customer_id', 'phone_no',
↪'year'])

```

```

[34]: running_model(X_train, X_test, y_train, y_test)

```

	precision	recall	f1-score	support
0.0	0.90	0.98	0.94	513
1.0	0.47	0.13	0.20	63
accuracy			0.89	576
macro avg	0.69	0.55	0.57	576
weighted avg	0.85	0.89	0.86	576

The area under the curve is: 0.55

First we executed the `prep_model` function to define our training and test sets as per the function. Then we executed `running_model` function using the training and testing sets.

The precision, recall and f1-score for our prediction is not that good as per the model. Infact the accuracy is very good, so that means our model is predicting properly which customers are not getting churned, but it is not predicting that accurately which customers are getting churned. Also, the area under the curve is 0.55 which is not that good as the worst model has 0.5 AUC.

```
[35]: #Confusion matrix

def conf_mat(y_test, y_pred):
    from sklearn.metrics import confusion_matrix

    confusion_matrix = confusion_matrix(y_test, y_pred)
    print(confusion_matrix)

    tn, fp, fn, tp = confusion_matrix.ravel()
    print('TN: %0.2f'%tn)
    print('FP: %0.2f'%fp)
    print('FN: %0.2f'%fn)
    print('TP: %0.2f'%tp)
```

```
[40]: #ROC Curve

def roc_cur(logreg, X_test, y_test):
    import matplotlib.pyplot as plt
    from sklearn.metrics import roc_auc_score
    from sklearn.metrics import roc_curve

    logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))

    fpr, tpr, thresholds = roc_curve(y_test, logreg.predict(X_test))

    #Creating Graph
    plt.figure()
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
```



```

#Plotting the worst ROC model
plt.plot([0, 1], [0, 1], 'b--')

#Plotting the logistic regression which we have built
plt.plot(fpr, tpr, color='darkorange', label='Logistic Regression (area =_
→%0.2f)'%logit_roc_auc)

#Adding labels and titles
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.savefig('LogROC')
plt.show()

```

```
[38]: conf_mat(y_test, y_pred)
```


```

[[504   9]
 [ 55   8]]
TN: 504.00
FP: 55.00
FN: 55.00
TP: 8.00

```

Our TN = 504, TP = 8, this is the reason our model is not good.

```
[41]: roc_cur(logreg, X_test, y_test)
```



This curve shows that our logistic regression model is not very far from the worst model, hence our model choice is not that good.

## 1.7 Dealing with Class Imbalance

```
[42]: df['churn'].describe()
```

```
[42]: count    1918.000000
      mean      0.131908
      std       0.338479
      min       0.000000
      25%       0.000000
      50%       0.000000
      75%       0.000000
      max       1.000000
      Name: churn, dtype: float64
```

In this project we're trying to predict how many customers are getting churned from the platform. So, the churn attribute is our dependent variable. When we describe this attribute we can see that our mean is 0.13, i.e. 13% which says that amount of people getting churned for the year 2015 are only 13%. This brings in the issue of class imbalance and for the same reason our model doesn't work well on this dataset.

## 1. Dealing with class imbalance using `class_weight = balanced` in `LogisticRegression` function of `sklearn`

```
[43]: def run_model_bweights(X_train, X_test, y_train, y_test):

    from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import roc_auc_score, classification_report

    #Defining the logistic model globally to use it outside the function
    global logreg

    #Fitting the logistic regression model from sklearn
    logreg = LogisticRegression(random_state = 13, class_weight = 'balanced')
    ↪ #here we're adding an extra attribute for balanced class weights
    logreg.fit(X_train, y_train)

    #Predicting y values
    global y_pred #Defining globally to use outside the function

    y_pred = logreg.predict(X_test)

    logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))

    print(classification_report(y_test, y_pred))
    print('The area under the curve is: %0.2f'%logit_roc_auc)
```

Using `class_weight = 'balanced'` the model either decreases the weight of the majority class or increases the weight of the minority class.

```
[44]: run_model_bweights(X_train, X_test, y_train, y_test)
```

	precision	recall	f1-score	support
0.0	0.96	0.74	0.84	513
1.0	0.27	0.78	0.40	63
accuracy			0.74	576
macro avg	0.62	0.76	0.62	576
weighted avg	0.89	0.74	0.79	576

The area under the curve is: 0.76

Now, we can see that our results are more better. We have lower precision and better recall, as a result, the AUC is 76%. This is a better model than the one which we executed above. Our F1-Score increased, but the accuracy had a dip of few percentages.

## 2. Dealing with class imbalance using `class_weight = w` as argument in `LogisticRegression` function of `sklearn` to pass our own tuned weights

```
[45]: def run_model_aweights(X_train, X_test, y_train, y_test, w):

    from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import roc_auc_score, classification_report

    #Defining the logistic model globally to use it outside the function
    global logreg

    #Fitting the logistic regression model from sklearn

    #here we're adding an extra attribute for our own passed class weights
    logreg = LogisticRegression(random_state = 13, class_weight = w)
    logreg.fit(X_train, y_train)

    #Predicting y values
    global y_pred #Defining globally to use outside the function

    y_pred = logreg.predict(X_test)

    logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))

    print(classification_report(y_test, y_pred))
    print('The area under the curve is: %0.2f'%logit_roc_auc)
```

```
[51]: run_model_aweights(X_train, X_test, y_train, y_test, {0:90, 1:10})
```

	precision	recall	f1-score	support
0.0	0.89	1.00	0.94	513
1.0	1.00	0.02	0.03	63
accuracy			0.89	576
macro avg	0.95	0.51	0.49	576
weighted avg	0.90	0.89	0.84	576

The area under the curve is: 0.51

Now, we can see our precision is really good, recall, F1-score and AUC is not good, but accuracy is better than earlier one.

### 3. Dealing with class imbalance using bootstraps for minority/majority class or re-sampling from majority/minority class

```
[53]: def adjust_imbalance(X_train, y_train, class_col):

    from sklearn.utils import resample #used to resample array or matrix in
    →consistent way
    import pandas as pd
```

```

X = pd.concat([X_train, y_train], axis=1)

#seperating 2 classes
class0 = X[X[class_col]==0]
class1 = X[X[class_col]==1]

#Case 1 - bootstrapping from minority class - used to convert the minority
→class into the size of majority class
if len(class1) < len(class0):
    resampled = resample(class1, replace=True, n_samples=len(class0),
→random_state=10)
    resampled_df = pd.concat([resampled, class0])

#Case 2 - Resampling from the majority class - used to cut cases from
→majority class until it has size of minority class
else:
    resampled = resample(class1, replace=False, n_samples=len(class0),
→random_state=10)
    resampled_df = pd.concat([resampled, class0])

return resampled_df

```

Above function will generate a resampled dataframe and we can run the model on resampled dataframe.

```

[54]: resampled_df = adjust_imbalance(X_train, y_train, class_col = 'churn')

[56]: prep_model(resampled_df, class_col = 'churn', cols_to_exclude = ['customer_id',
→'phone_no', 'year'])
running_model(X_train, X_test, y_train, y_test)

```

	precision	recall	f1-score	support
0.0	0.67	0.75	0.71	339
1.0	0.73	0.65	0.69	353
accuracy			0.70	692
macro avg	0.70	0.70	0.70	692
weighted avg	0.70	0.70	0.70	692

The area under the curve is: 0.70

This method gives us less accuracy but better AUC. The precision, recall and f1-score is also good here.

#### 4. Dealing with class imbalance using smote

```
[71]: #Synthetic Minority Optimization Technique. Generated new instances from
      ↪existing minority cases that supply as input.

def prep_model_smote(df, class_col, cols_to_exclude):

    from sklearn.model_selection import train_test_split
    import numpy as np
    from imblearn.over_sampling import SMOTE

    #cleaning the dataframe for logistic regression with the columns which
    ↪we're not using, i.e. phone no, customer_id, & year
    cols = df.select_dtypes(include=np.number).columns.tolist()
    X = df[cols]
    X = X[X.columns.difference(cols_to_exclude)]
    X = X[X.columns.difference([class_col])]

    y = df[class_col]

    #Declaring globally so that we can call this variables outside this
    ↪function also.
    global X_train, X_test, y_train, y_test

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
    ↪random_state = 0)

    sm = SMOTE(random_state=0, sampling_strategy=1.0)
    X_train, y_train = sm.fit_sample(X_train, y_train)

[72]: prep_model_smote(df, class_col = 'churn', cols_to_exclude = ['customer_id',
    ↪'phone_no', 'year'])
    running_model(X_train, X_test, y_train, y_test)
```

	precision	recall	f1-score	support
0.0	0.96	0.72	0.82	513
1.0	0.25	0.76	0.37	63
accuracy			0.72	576
macro avg	0.60	0.74	0.60	576
weighted avg	0.88	0.72	0.77	576

The area under the curve is: 0.74

Using this technique we get area under the curve of 76%, but with lower precision and higher recall.

## 1.8 Feature Selection

We will try to select best features for our model to increase the accuracy of our model with efficiency. This is beneficial when datasets are large and we can build the best model with less features.

**1. Feature Selection using Variance Threshold selection** We only get features which pass a certain variance threshold

```
[79]: class_col = 'churn'
      cols_to_exclude = ['customer_id', 'phone_no', 'year']

      def var_threshold_selection(df, cols_to_exclude, class_col, threshold):

          from sklearn.feature_selection import VarianceThreshold
          import numpy as np
          from sklearn import preprocessing

          #cleaning the dataframe for logistic regression with the columns which
          ↪we're not using, i.e. phone no, customer_id, & year
          cols = df.select_dtypes(include=np.number).columns.tolist()
          X = df[cols]
          X = X[X.columns.difference(cols_to_exclude)]
          X = X[X.columns.difference([class_col])]

          #Scaling Variables
          scaler = preprocessing.StandardScaler().fit(X)
          X_scaled = scaler.transform(X)
          var_thr = VarianceThreshold(threshold = threshold) #removing both constant
          ↪and quasi-constant
          var_thr.fit(X_scaled)
          var_thr.get_support()

          global selected_cols
          selected_cols = X.columns[var_thr.get_support()]

          print('The selected features are: ')
          print(list(selected_cols))
```

```
[83]: var_threshold_selection(df, cols_to_exclude=['customer_id', 'phone_no',
          ↪'year'], class_col = 'churn', threshold = 1)
```

The selected features are:

['maximum\_daily\_mins', 'maximum\_days\_inactive', 'weekly\_mins\_watched']

```
[84]: prep_model(resampled_df, class_col = 'churn', cols_to_exclude=['customer_id',
          ↪'phone_no', 'year',
          ↪'gender', 'age',
          ↪'no_of_days_subscribed',
```

```

                                'multi_screen',
↪ 'mail_subscribed', 'minimum_daily_mins',
                                ↵
↪ 'weekly_max_night_mins', 'videos_watched',
                                ↵
↪ 'customer_support_calls', 'churn', 'gender_code',
                                ↵
↪ 'multi_screen_code', 'mail_subscribed_code'])
running_model(X_train, X_test, y_train, y_test)

```

	precision	recall	f1-score	support
0.0	0.57	0.62	0.59	339
1.0	0.60	0.55	0.57	353
accuracy			0.58	692
macro avg	0.58	0.58	0.58	692
weighted avg	0.59	0.58	0.58	692

The area under the curve is: 0.58

This is not that good technique of feature selection as the threshold is not tuned properly and many columns are getting dropped. There is better technique than this which is called as RFE. Above model is similar to the previous one and did not brought a lot of improvement.

**2. Feature Selection using Recursive Feature Eliminations (RFE)** This method executes several models with different features and eliminates them one by one to see the optimal amount of features.

```

[89]: def rfe_selection(df, cols_to_exclude, class_col, model):

    import warnings
    warnings.filterwarnings('ignore')

    from sklearn.feature_selection import RFE
    import numpy as np

    #cleaning the dataframe for logistic regression with the columns which
↪ we're not using, i.e. phone no, customer_id, & year
    cols = df.select_dtypes(include=np.number).columns.tolist()
    X = df[cols]
    X = X[X.columns.difference(cols_to_exclude)]
    X = X[X.columns.difference([class_col])]
    y = df[class_col]

    rfe = RFE(model)
    rfe = rfe.fit(X, y)

```



```

global selected_cols
selected_cols = X.columns[rfe.support_]

print('The selected features are: ')
print(list(selected_cols))

```

```

[90]: rfe_selection(df, cols_to_exclude=['customer_id', 'phone_no', 'year'],
    ↪class_col = 'churn', model=logreg)

```

The selected features are:

```

['customer_support_calls', 'gender_code', 'mail_subscribed_code',
'maximum_days_inactive', 'minimum_daily_mins', 'multi_screen_code']

```

```

[92]: prep_model(resampled_df, class_col = 'churn', cols_to_exclude=['customer_id',
    ↪'phone_no', 'year',
    ↪'gender', 'age',
    ↪'no_of_days_subscribed',
    ↪'multi_screen',
    ↪'mail_subscribed',
    ↪'weekly_max_night_mins', 'gender_code',
    ↪'multi_screen_code', 'mail_subscribed_code'])
running_model(X_train, X_test, y_train, y_test)

```

	precision	recall	f1-score	support
0.0	0.69	0.71	0.70	339
1.0	0.71	0.69	0.70	353
accuracy			0.70	692
macro avg	0.70	0.70	0.70	692
weighted avg	0.70	0.70	0.70	692

The area under the curve is: 0.70

This model gives more better result than the earlier one, these results are very much similar to the one where we used resampled\_df. It gives the similar results with very less features. So, it's a better model.

## 1.9 Saving and Running the Model

```

[96]: import pickle
pickle.dump(logreg, open('model1.pk1', 'wb')) #use to dump the logreg model
    ↪into the file name mentioned in the command

```

```
[94]: model = pickle.load(open('model1.pk1', 'rb'))
```

```
[97]: model.predict(X_test) #Sample Prediction
```

```
[97]: array([1., 0., 0., 1., 1., 0., 1., 1., 0., 0., 1., 1., 1., 1., 0., 0., 1.,  
1., 0., 1., 1., 1., 1., 0., 0., 0., 0., 1., 1., 0., 0., 1., 1., 0.,  
1., 1., 1., 0., 0., 1., 0., 1., 0., 1., 1., 1., 0., 1., 1., 1., 0.,  
1., 0., 1., 1., 1., 1., 0., 0., 0., 1., 0., 0., 0., 1., 0., 0., 0.,  
1., 1., 0., 1., 1., 0., 0., 0., 1., 0., 1., 0., 0., 1., 0., 1., 1.,  
1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 0., 1., 0., 1., 0., 0., 0.,  
1., 1., 1., 1., 0., 0., 1., 1., 0., 1., 0., 1., 1., 1., 1., 0., 1.,  
1., 0., 1., 1., 0., 1., 0., 0., 0., 0., 1., 0., 1., 1., 0., 0., 1.,  
1., 1., 1., 0., 1., 1., 0., 1., 0., 0., 1., 1., 0., 1., 0., 1., 0.,  
0., 0., 1., 0., 1., 1., 0., 0., 1., 0., 1., 0., 1., 0., 0., 1., 0.,  
1., 0., 1., 1., 0., 1., 1., 1., 0., 1., 0., 0., 0., 0., 0., 0., 0.,  
0., 0., 1., 0., 1., 0., 1., 1., 1., 0., 1., 1., 0., 0., 0., 0., 1.,  
1., 1., 1., 0., 0., 0., 0., 0., 0., 1., 0., 1., 1., 0., 0., 0., 0.,  
1., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 1., 1., 0.,  
0., 1., 0., 1., 1., 1., 1., 0., 0., 0., 1., 1., 1., 0., 1., 0., 0.,  
0., 0., 1., 1., 0., 1., 0., 1., 1., 0., 1., 1., 0., 0., 1., 0., 0.,  
1., 1., 0., 1., 1., 1., 1., 1., 1., 0., 0., 1., 0., 0., 0., 0., 0.,  
1., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 1., 1., 0.,  
1., 1., 0., 1., 1., 0., 1., 1., 1., 0., 0., 1., 1., 1., 0., 0., 0.,  
0., 0., 1., 1., 0., 0., 1., 0., 0., 1., 0., 1., 1., 1., 0., 0., 0.,  
0., 1., 0., 1., 1., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 1., 1.,  
1., 1., 1., 1., 1., 1., 0., 0., 0., 0., 1., 1., 0., 0., 0., 0., 1.,  
0., 1., 0., 1., 0., 1., 0., 0., 1., 0., 1., 0., 0., 0., 0., 1., 1.,  
1., 1., 1., 0., 0., 1., 0., 0., 0., 0., 0., 1., 0., 0., 1., 0., 1.,  
1., 0., 0., 1., 1., 0., 0., 1., 0., 0., 1., 0., 1., 0., 1., 0., 1.,  
1., 1., 0., 1., 0., 0., 1., 0., 0., 1., 0., 1., 1., 1., 0., 0., 0.,  
1., 1., 0., 1., 0., 0., 0., 0., 0., 0., 1., 1., 0., 1., 0., 0., 1.,  
1., 0., 0., 0., 0., 0., 0., 0., 0., 1., 1., 0., 1., 0., 0., 0., 1.,  
0., 1., 1., 1., 0., 0., 1., 0., 0., 1., 0., 1., 1., 1., 0., 0.,  
1., 0., 1., 0., 0., 0., 0., 0., 0., 1., 1., 1., 0., 1., 1., 1., 0.,  
0., 0., 0., 0., 0., 1., 1., 0., 0., 1., 1., 1.]
```